

# U1M5.LW.Access and Join Methods Part 2

## Report

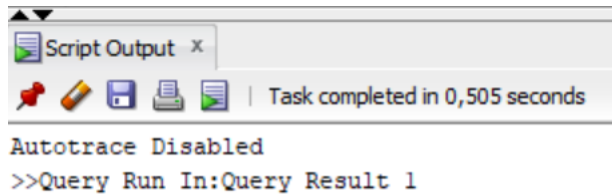
Mastykina Elizaweta

## 1. Auto Trace & Explain Plan

### 1.1. Task 1: Auto Trace configuration training

**SELECT emp.\* FROM emp;**

1. set autotrace off



2. set autotrace on

Id	Operation	Name	E-Rows
0	SELECT STATEMENT		
1	TABLE ACCESS FULL	EMP	60720

PLAN\_TABLE\_OUTPUT

Note

-----

- Warning: basic plan statistics not available. These are only collected when:  
\* hint 'gather\_plan\_statistics' is used for the statement or  
\* parameter 'statistics\_level' is set to 'ALL', at session or system level

Statistics

-----

2	CPU used by this session
5	CPU used when call started
4	DB time
43	Requests to/from client
43	SQL*Net roundtrips to/from client
601	bytes received via SQL*Net from client
215582	bytes sent via SQL*Net to client
2	calls to get snapshot scn: kcmgss
4	calls to kcmgcs
19	consistent gets
19	consistent gets from cache
19	consistent gets pin
19	consistent gets pin (fastpath)
2	cursor authentications
2	execute count
622592	logical read bytes from cache
17	no work - consistent read gets
70	non-idle wait count
2	opened cursors cumulative
1	opened cursors current
2	parse count (total)
7	process last non-idle time
19	session logical reads
1	sorts (memory)
1581	sorts (rows)
17	table scan blocks gotten
12529	table scan disk non-IMC rows gotten
12529	table scan rows gotten
1	table scans (short tables)
45	user calls

### 3. set autotrace traceonly

Not supported on SQL Developer.

### 4. set autotrace traceonly explain

Not supported on SQL Developer

### 5. set autotrace traceonly statistics

Not supported on SQL Developer

### 6. set autotrace traceonly explain statistics

Not supported on SQL Developer.

### 7. set autotrace on explain

```
PLAN_TABLE_OUTPUT
-----
SQL_ID a6yp0kv8klg9s, child number 0
-----
SELECT emp.* FROM emp

Plan hash value: 3956160932

-----
| Id | Operation          | Name | E-Rows |
-----
| 0  | SELECT STATEMENT   |      |        |
| 1  | TABLE ACCESS FULL| EMP  | 60720  |
-----
PLAN_TABLE_OUTPUT
-----
```

### 8. set autotrace on statistics

```
Statistics
-----
      2 CPU used by this session
      2 CPU used when call started
     43 Requests to/from client
     43 SQL*Net roundtrips to/from client
    602 bytes received via SQL*Net from client
  212394 bytes sent via SQL*Net to client
      3 calls to get snapshot scn: kcmgss
      4 calls to kcmgcs
     19 consistent gets
     19 consistent gets from cache
     19 consistent gets pin
     19 consistent gets pin (fastpath)
      1 enqueue releases
      1 enqueue requests
      2 execute count
  622592 logical read bytes from cache
     17 no work - consistent read gets
     57 non-idle wait count
      2 opened cursors cumulative
      1 opened cursors current
      1 parse count (hard)
      2 parse count (total)
      8 process last non-idle time
      1 recursive calls
     19 session logical reads
      1 sorts (memory)
   1581 sorts (rows)
     17 table scan blocks gotten
  12529 table scan disk non-IMC rows gotten
  12529 table scan rows gotten
      1 table scans (short tables)
     45 user calls
```

9. set autotrace on explain statistics

The same as set autotrace on.

10. set autotrace off explain

Currently isn't used.

11. set autotrace off statistics

Currently isn't used.

12. set autotrace off explain statistics

Currently isn't used.

Auto Trace Configuration Options	Results and Description
SET autotrace off	no AUTOTRACE report is generated, this is the default
SET AUTOTRACE ON	the AUTOTRACE report includes both the optimizer execution path and the SQL statement execution statistics
SET AUTOTRACE TRACEONLY	the AUTOTRACE report includes both the optimizer execution path and the SQL statement execution statistics, but suppresses the printing of the user's query output! not supported on SQL Developer
SET AUTOTRACE ON EXPLAIN	the AUTOTRACE report shows only the optimizer execution path
SET AUTOTRACE ON STATISTICS	the AUTOTRACE report shows only the SQL statement execution statistics
SET AUTOTRACE ON EXPLAIN STATISTICS	the AUTOTRACE report includes both the optimizer execution path and the SQL statement execution statistics
SET AUTOTRACE TRACEONLY EXPLAIN	the AUTOTRACE report includes the optimizer execution path, but suppresses the printing of the user's query output! currently not used
SET AUTOTRACE TRACEONLY STATISTICS	the AUTOTRACE report includes the SQL statement execution statistics, but suppresses the printing of the user's query output! not supported on SQL Developer
SET AUTOTRACE TRACEONLY EXPLAIN STATISTICS	the AUTOTRACE report includes both the optimizer execution path and the SQL statement execution statistics, but suppresses the printing of the user's query output! not supported on SQL Developer
SET AUTOTRACE OFF EXPLAIN	Is not used
SET AUTOTRACE OFF STATISTICS	Is not used
SET AUTOTRACE OFF EXPLAIN STATISTICS	Is not used

## 2. Join Methods

0,1879999 seconds

Worksheet Query Builder

```
SELECT *
FROM scott.emp e, scott.dept d
WHERE e.deptno = d.deptno
AND d.deptno = 10
```

Script Output x

Task completed in 0,188 seconds

EMPNO	ENAME	JOB	MGR	HIREDATE	SAL	COMM	DEPTNO	DEPTNO DNAME	LOC
7782	CLARK	MANAGER	7839	09.06.81	2450		10	10 ACCOUNTING	NEW YORK
7839	KING	PRESIDENT		17.11.81	5000		10	10 ACCOUNTING	NEW YORK
7934	MILLER	CLERK	7782	23.01.82	1300		10	10 ACCOUNTING	NEW YORK

SQL Worksheet History

48,45000076 seconds

Worksheet Query Builder

```
FROM scott.emp e, scott.dept d
WHERE e.deptno = d.deptno
AND d.deptno = 10
```

Script Output x

Task completed in 48,45 seconds

Autotrace Enabled  
Displays the execution plan only.

EMPNO	ENAME	JOB	MGR	HIREDATE	SAL	COMM	DEPTNO	DEPTNO DNAME	LOC
7782	CLARK	MANAGER	7839	09.06.81	2450		10	10 ACCOUNTING	NEW YORK
7839	KING	PRESIDENT		17.11.81	5000		10	10 ACCOUNTING	NEW YORK
7934	MILLER	CLERK	7782	23.01.82	1300		10	10 ACCOUNTING	NEW YORK

PLAN\_TABLE\_OUTPUT

SQL\_ID 287skc8v2zgg0, child number 0

```
SELECT * FROM scott.emp e, scott.dept d WHERE e.deptno = d.deptno AND d.deptno = 10
```

Plan hash value: 568005898

Id	Operation	Name	E-Rows
0	SELECT STATEMENT		

PLAN\_TABLE\_OUTPUT

Id	Operation	Name	E-Rows
1	NESTED LOOPS		5
2	TABLE ACCESS BY INDEX ROWID	DEPT	1
* 3	INDEX UNIQUE SCAN	PK_DEPT	1
* 4	TABLE ACCESS FULL	EMP	5

Predicate Information (identified by operation id):

```
3 - access("D"."DEPTNO"=10)
4 - filter("E"."DEPTNO"=10)
```

PLAN\_TABLE\_OUTPUT

Script Output x Explain Plan x				
SQL   0,055 seconds				
OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			3	3
NESTED LOOPS			3	3
TABLE ACCESS	DEPT	BY INDEX ROWID	1	2
INDEX	SYS_C007178	UNIQUE SCAN	1	1
Access Predicates	D.DEPTNO=10			
TABLE ACCESS	EMP	CLUSTER	3	1
INDEX	IDXCL_EMP_DEPT	UNIQUE SCAN	1	0
Access Predicates	E.DEPTNO=10			

The conception of work of Nested Loops Join is the following: Oracle takes the first value from the first table("outer" table is chosen by default) and compare to each value in "inner" table, in the search of match.

When all the values in the 'inner' table are checked Oracle takes next value in the first table and repeat this process. The same for all the rest values in the first table.

It is important to notice, that such method of join is the most ineffective. But we have sum methods to optimize it.

## 2.2. Task 3: Sort-Merge Joins

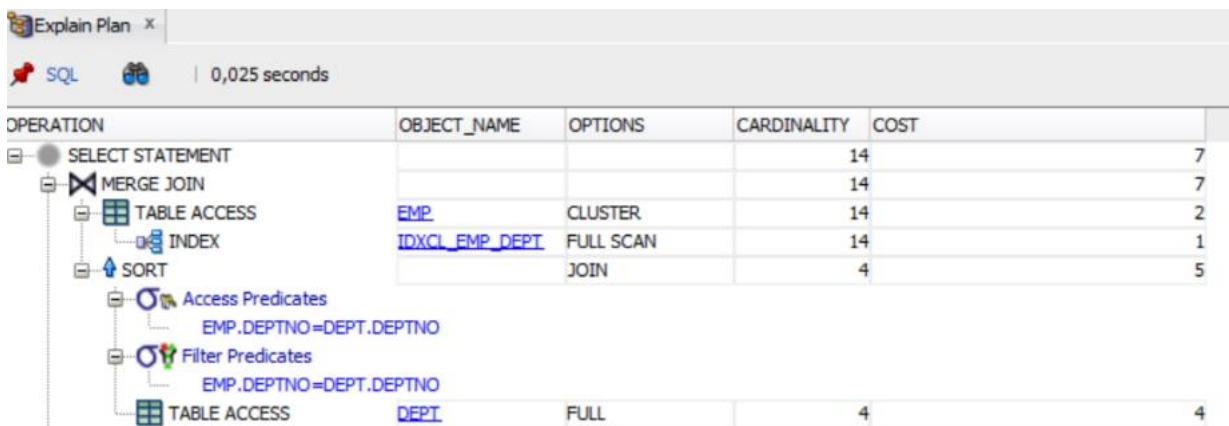
```

SET AUTOTRACE ON EXPLAIN;
SELECT /*+ USE_MERGE(dept emp) */
FROM dept, emp
WHERE emp.deptno= dept.deptno;
SELECT * FROM TABLE(dbms_xplan.display );

```

Script Output x Explain Plan x Query Result x Query Result 1 x										
SQL   All Rows Fetched: 14 in 0,02 seconds										
DEPTNO	DNAME	LOC	EMPNO	ENAME	JOB	MGR	HIREDATE	SAL	COMM	DEPTNO_1
1	10 ACCOUNTING	NEW YORK	7782	CLARK	MANAGER	7839	09.06.81	2450	(null)	10
2	10 ACCOUNTING	NEW YORK	7839	KING	PRESIDENT	(null)	17.11.81	5000	(null)	10
3	10 ACCOUNTING	NEW YORK	7934	MILLER	CLERK	7782	23.01.82	1300	(null)	10
4	20 RESEARCH	DALLAS	7369	SMITH	CLERK	7902	17.12.80	800	(null)	20
5	20 RESEARCH	DALLAS	7566	JONES	MANAGER	7839	02.04.81	2975	(null)	20
6	20 RESEARCH	DALLAS	7788	SCOTT	ANALYST	7566	13.07.87	3000	(null)	20
7	20 RESEARCH	DALLAS	7876	ADAMS	CLERK	7788	13.07.87	1100	(null)	20
8	20 RESEARCH	DALLAS	7902	FORD	ANALYST	7566	03.12.81	3000	(null)	20
9	30 SALES	CHICAGO	7499	ALLEN	SALESMAN	7698	20.02.81	1600	300	30
10	30 SALES	CHICAGO	7521	WARD	SALESMAN	7698	22.02.81	1250	500	30
11	30 SALES	CHICAGO	7654	MARTIN	SALESMAN	7698	28.09.81	1250	1400	30
12	30 SALES	CHICAGO	7698	BLAKE	MANAGER	7839	01.05.81	2850	(null)	30
13	30 SALES	CHICAGO	7844	TURNER	SALESMAN	7698	13.07.87	1500	0	30
14	30 SALES	CHICAGO	7900	JAMES	CLERK	7698	03.12.81	950	(null)	30

PLAN_TABLE_OUTPUT							
1	Plan hash value: 3809461229						
2							
3	-----						
4	Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
5	-----						
6	0	SELECT STATEMENT		14	1638	7 (15)	00:00:01
7	1	MERGE JOIN		14	1638	7 (15)	00:00:01
8	2	TABLE ACCESS CLUSTER	EMP	14	1218	2 (0)	00:00:01
9	3	INDEX FULL SCAN	IDXCL_EMP_DEPT	14		1 (0)	00:00:01
10	* 4	SORT JOIN		4	120	5 (20)	00:00:01
11	5	TABLE ACCESS FULL	DEPT	4	120	4 (0)	00:00:01



Sort-merge joins read two tables independently and sort rows (according to WHERE statment) from each table in the order of connection key and then joins sorted rows. In case we have already sorted list we don't apply sort again. When we work with big amount of source, which don't fit memory, sort uses Operation memory.

## 2.3. Task 4: Hash Joins

```

SET AUTOTRACE ON EXPLAIN;
SELECT /*+ USE_HASH(dept emp) */ *
FROM dept, emp
WHERE emp.deptno= dept.deptno;
SELECT * FROM TABLE(dbms_xplan.display );

```



DEPTNO	DNAME	LOC	EMPNO	ENAME	JOB	MGR	HIREDATE	SAL	COMM	DEPTNO_1
1	10 ACCOUNTING	NEW YORK	7782 CLARK	MANAGER	7839 09.06.81	2450	(null)	10		
2	10 ACCOUNTING	NEW YORK	7839 KING	PRESIDENT	(null) 17.11.81	5000	(null)	10		
3	10 ACCOUNTING	NEW YORK	7934 MILLER	CLERK	7782 23.01.82	1300	(null)	10		
4	20 RESEARCH	DALLAS	7369 SMITH	CLERK	7902 17.12.80	800	(null)	20		
5	20 RESEARCH	DALLAS	7566 JONES	MANAGER	7839 02.04.81	2975	(null)	20		
6	20 RESEARCH	DALLAS	7788 SCOTT	ANALYST	7566 13.07.87	3000	(null)	20		
7	20 RESEARCH	DALLAS	7876 ADAMS	CLERK	7788 13.07.87	1100	(null)	20		
8	20 RESEARCH	DALLAS	7902 FORD	ANALYST	7566 03.12.81	3000	(null)	20		
9	30 SALES	CHICAGO	7499 ALLEN	SALESMAN	7698 20.02.81	1600	300	30		
0	30 SALES	CHICAGO	7521 WARD	SALESMAN	7698 22.02.81	1250	500	30		
1	30 SALES	CHICAGO	7654 MARTIN	SALESMAN	7698 28.09.81	1250	1400	30		
2	30 SALES	CHICAGO	7698 BLAKE	MANAGER	7839 01.05.81	2850	(null)	30		
3	30 SALES	CHICAGO	7844 TURNER	SALESMAN	7698 13.07.87	1500	0	30		
4	30 SALES	CHICAGO	7900 JAMES	CLERK	7698 03.12.81	950	(null)	30		

PLAN_TABLE_OUTPUT							
1	Plan hash value: 2553673404						
2							
3	-----						
Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time	
4	-----						
5							
6	0	SELECT STATEMENT	14	1638	6 (0)	00:00:01	
7	* 1	HASH JOIN	14	1638	6 (0)	00:00:01	
8	2	TABLE ACCESS FULL	4	120	4 (0)	00:00:01	
9	3	TABLE ACCESS CLUSTER	14	1218	2 (0)	00:00:01	
10	4	INDEX FULL SCAN	14		1 (0)	00:00:01	

Explain Plan				
SQL   0,021 seconds				
OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			14	6
HASH JOIN			14	6
Access Predicates				
EMP.DEPTNO=DEPT.DEPTNO				
TABLE ACCESS	DEPT	FULL	4	4
TABLE ACCESS	EMP	CLUSTER	14	2
INDEX	IDXCL_EMP_DEPT	FULL SCAN	14	1
Other XML				

Hash joins reads two tables and applies WHERE condition. Base on the table statistic, the table which returns the least number of rows hashing in the memory. This hash table includes all table row data and is loaded into hash blocks based on a randomization function that converts the join key to a hash value. As long as there is enough memory, this hash table will remain in memory. However, if there is not enough available memory, the hash table can be written to temporary disk space. The next step is to read another larger table and apply a hash function to the join key column. This hash value is then used to check the smaller in-memory table for the corresponding hash bucket



containing the row data for the first table. Each bin has a list (represented by a bitmap) of the rows in that bin. This list is checked against the test string. If a match is found, the string is returned, otherwise it is discarded. The large line is read only once and each line is checked for a match. This is different from Nested Loops Joins, where the inner table is read multiple times. So actually in this case the big table is the leading table since they are only read once and the smaller hash table is read many times.

## 2.4. Task 5: Cartesian Joins

```
SET AUTOTRACE ON EXPLAIN;
SELECT * FROM dept, emp;
SELECT * FROM TABLE(dbms_xplan.display );
```

	DEPTNO	DNAME	LOC	EMPNO	ENAME	JOB	MGR	HIREDATE	SAL	COMM	DEPTNO_1
1	10	ACCOUNTING	NEW YORK	7369	SMITH	CLERK	7902	17.12.80	800	(null)	20
2	10	ACCOUNTING	NEW YORK	7499	ALLEN	SALESMAN	7698	20.02.81	1600	300	30
3	10	ACCOUNTING	NEW YORK	7521	WARD	SALESMAN	7698	22.02.81	1250	500	30
4	10	ACCOUNTING	NEW YORK	7566	JONES	MANAGER	7839	02.04.81	2975	(null)	20
5	10	ACCOUNTING	NEW YORK	7654	MARTIN	SALESMAN	7698	28.09.81	1250	1400	30
6	10	ACCOUNTING	NEW YORK	7698	BLAKE	MANAGER	7839	01.05.81	2850	(null)	30
7	10	ACCOUNTING	NEW YORK	7782	CLARK	MANAGER	7839	09.06.81	2450	(null)	10
8	10	ACCOUNTING	NEW YORK	7788	SCOTT	ANALYST	7566	13.07.87	3000	(null)	20
9	10	ACCOUNTING	NEW YORK	7839	KING	PRESIDENT	(null)	17.11.81	5000	(null)	10
10	10	ACCOUNTING	NEW YORK	7844	TURNER	SALESMAN	7698	13.07.87	1500	0	30
11	10	ACCOUNTING	NEW YORK	7876	ADAMS	CLERK	7788	13.07.87	1100	(null)	20
12	10	ACCOUNTING	NEW YORK	7900	JAMES	CLERK	7698	03.12.81	950	(null)	30
13	10	ACCOUNTING	NEW YORK	7902	FORD	ANALYST	7566	03.12.81	3000	(null)	20
14	10	ACCOUNTING	NEW YORK	7934	MILLER	CLERK	7782	23.01.82	1300	(null)	10
15	20	RESEARCH	DALLAS	7369	SMITH	CLERK	7902	17.12.80	800	(null)	20
16	20	RESEARCH	DALLAS	7499	ALLEN	SALESMAN	7698	20.02.81	1600	300	30
17	20	RESEARCH	DALLAS	7521	WARD	SALESMAN	7698	22.02.81	1250	500	30

PLAN_TABLE_OUTPUT									
1	Plan hash value: 2034389985								
2									
3	-----								
4	Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time		
5	-----								
6	0	SELECT STATEMENT		56	6552	12 (0)	00:00:01		
7	1	MERGE JOIN CARTESIAN		56	6552	12 (0)	00:00:01		
8	2	TABLE ACCESS FULL	DEPT	4	120	4 (0)	00:00:01		
9	3	BUFFER SORT		14	1218	8 (0)	00:00:01		
10	4	TABLE ACCESS FULL	EMP	14	1218	2 (0)	00:00:01		
11									

Explain Plan x				
SQL 0,016 seconds				
OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			56	12
MERGE JOIN		CARTESIAN	56	12
TABLE ACCESS	DEPT	FULL	4	4
BUFFER		SORT	14	8
TABLE ACCESS	EMP	FULL	14	2

## 2.5. Task 6: Left/Right Outer Joins

Worksheet

Query Builder

```
SELECT * FROM scott.emp e LEFT OUTER JOIN scott.dept d
ON e.deptno = d.deptno;
```

Script Output x

Task completed in 0,183 seconds

EMPNO	ENAME	JOB	MGR	HIREDATE	SAL	COMM	DEPTNO	DEPTNO	DNAME	LOC
7782	CLARK	MANAGER	7839	09.06.81	2450		10	10	ACCOUNTING	NEW YORK
7839	KING	PRESIDENT		17.11.81	5000		10	10	ACCOUNTING	NEW YORK
7934	MILLER	CLERK	7782	23.01.82	1300		10	10	ACCOUNTING	NEW YORK
7369	SMITH	CLERK	7902	17.12.80	800		20	20	RESEARCH	DALLAS
7566	JONES	MANAGER	7839	02.04.81	2975		20	20	RESEARCH	DALLAS
7788	SCOTT	ANALYST	7566	19.04.87	3000		20	20	RESEARCH	DALLAS
7876	ADAMS	CLERK	7788	23.05.87	1100		20	20	RESEARCH	DALLAS
7902	FORD	ANALYST	7566	03.12.81	3000		20	20	RESEARCH	DALLAS
7499	ALLEN	SALESMAN	7698	20.02.81	1600	300	30	30	SALES	CHICAGO
7521	WARD	SALESMAN	7698	22.02.81	1250	500	30	30	SALES	CHICAGO
7654	MARTIN	SALESMAN	7698	28.09.81	1250	1400	30	30	SALES	CHICAGO
7698	BLAKE	MANAGER	7839	01.05.81	2850		30	30	SALES	CHICAGO
7844	TURNER	SALESMAN	7698	08.09.81	1500	0	30	30	SALES	CHICAGO
7900	JAMES	CLERK	7698	03.12.81	950		30	30	SALES	CHICAGO

14 rows selected.

Cartesian Joins occur when all rows in one table are joined with all rows in another table. Therefore, the total number of rows resulting from the join is equal to the number of rows from one table (A) times the number of rows from the other table (B), so  $A \times B$  = total number of rows in the result set. Cartesian joins often occur when a join condition is omitted or a specified join column is ignored, so the only possible operation is to simply join all rows from one table with all rows from another table.

SELECT * FROM scott.emp e LEFT OUTER JOIN scott.dept d ON e.deptno = d.deptno										
Plan hash value: 3387915970										
Id	Operation	Name	Starts	E-Rows	A-Rows	A-Time	Buffers	OMem	1Mem	Used-Mem
0	SELECT STATEMENT		1		14	00:00:00.01	8			
PLAN_TABLE_OUTPUT										
* 1	HASH JOIN OUTER		1	14	14	00:00:00.01	8	604K	604K	627K (0)
2	TABLE ACCESS FULL	EMP	1	14	14	00:00:00.01	4			
3	TABLE ACCESS FULL	DEPT	1	4	4	00:00:00.01	4			

Predicate Information (identified by operation id):				
1 - access("E"."DEPTNO"="D"."DEPTNO")				
OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			14	8
HASH JOIN		OUTER	14	8
Access Predicates E.DEPTNO=D.DEPTNO(+)				
NESTED LOOPS		OUTER	14	8
STATISTICS COLLECTOR				
TABLE ACCESS	SCOTT.EMP	FULL	14	4
TABLE ACCESS	SCOTT.DEPT	BY INDEX ROWID	1	4
INDEX	SCOTT.PK_DEPT	UNIQUE SCAN		
Access Predicates E.DEPTNO=D.DEPTNO(+)				
TABLE ACCESS	SCOTT.DEPT	FULL	4	4
Other XML				

2.6. Task 7: Full Outer Join

SQL Worksheet | History

8,53299999 seconds

Worksheet | Query Builder

```
SELECT * FROM scott.emp e FULL OUTER JOIN scott.dept d
ON e.deptno = d.deptno;
```

Script Output x

Task completed in 8,533 seconds

EMPNO	ENAME	JOB	MGR	HIREDATE	SAL	COMM	DEPTNO	DEPTNO	DNAME	LOC
7369	SMITH	CLERK	7902	17.12.80	800		20	20	RESEARCH	DALLAS
7499	ALLEN	SALESMAN	7698	20.02.81	1600	300	30	30	SALES	CHICAGO
7521	WARD	SALESMAN	7698	22.02.81	1250	500	30	30	SALES	CHICAGO
7566	JONES	MANAGER	7839	02.04.81	2975		20	20	RESEARCH	DALLAS
7654	MARTIN	SALESMAN	7698	28.09.81	1250	1400	30	30	SALES	CHICAGO
7698	BLAKE	MANAGER	7839	01.05.81	2850		30	30	SALES	CHICAGO
7782	CLARK	MANAGER	7839	09.06.81	2450		10	10	ACCOUNTING	NEW YORK
7788	SCOTT	ANALYST	7566	19.04.87	3000		20	20	RESEARCH	DALLAS
7839	KING	PRESIDENT		17.11.81	5000		10	10	ACCOUNTING	NEW YORK
7844	TURNER	SALESMAN	7698	08.09.81	1500	0	30	30	SALES	CHICAGO
7876	ADAMS	CLERK	7788	23.05.87	1100		20	20	RESEARCH	DALLAS
7900	JAMES	CLERK	7698	03.12.81	950		30	30	SALES	CHICAGO
7902	FORD	ANALYST	7566	03.12.81	3000		20	20	RESEARCH	DALLAS
7934	MILLER	CLERK	7782	23.01.82	1300		10	10	ACCOUNTING	NEW YORK
								40	OPERATIONS	BOSTON

SELECT \* FROM scott.emp e FULL OUTER JOIN scott.dept d ON e.deptno = d.deptno

Plan hash value: 51889263

Id	Operation	Name	Starts	E-Rows	A-Rows	A-Time	Buffers	OMem	1Mem	Used-Mem
0	SELECT STATEMENT		1		15	00:00:00.01	8			

PLAN\_TABLE\_OUTPUT

1	VIEW	VW_FOJ_0	1	15	15	00:00:00.01	8			
* 2	HASH JOIN FULL OUTER		1	15	15	00:00:00.01	8	801K	801K	902K (0)
3	TABLE ACCESS FULL	DEPT	1	4	4	00:00:00.01	4			
4	TABLE ACCESS FULL	EMP	1	14	14	00:00:00.01	4			

Predicate Information (identified by operation id):

2 - access("E"."DEPTNO"="D"."DEPTNO")

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			15	8
VIEW	<a href="#">SYS.VW_FOJ_0</a>		15	8
HASH JOIN		FULL OUTER	15	8
Access Predicates				
E.DEPTNO=D.DEPTNO				
TABLE ACCESS	<a href="#">SCOTT.DEPT</a>	FULL	4	4
TABLE ACCESS	<a href="#">SCOTT.EMP</a>	FULL	14	4
Other XML				

2.7. Task 8: Semi Joins

SQL Worksheet History

0,19 seconds

Worksheet Query Builder

```
SELECT /* using in */ deptno, dname FROM scott.dept d
WHERE deptno IN (select deptno from scott.emp e);
```

Script Output x

Task completed in 0,19 seconds

DEPTNO	DNAME
10	ACCOUNTING
20	RESEARCH
30	SALES

```
SELECT /* using in */ deptno, dname FROM scott.dept d WHERE deptno IN
(select deptno from scott.emp e)
```

Plan hash value: 1090737117

Id	Operation	Name	Starts	E-Rows	A-Rows	A-Time	Buffers	OMem	IMem	Used-Mem
0	SELECT STATEMENT		1		3	00:00:00.01	6			

PLAN\_TABLE\_OUTPUT

1	MERGE JOIN SEMI		1	3	3	00:00:00.01	6			
2	TABLE ACCESS BY INDEX ROWID	DEPT	1	4	4	00:00:00.01	2			
3	INDEX FULL SCAN	PK_DEPT	1	4	4	00:00:00.01	1			
* 4	SORT UNIQUE		4	14	3	00:00:00.01	4	2048	2048	2048 (0)
5	TABLE ACCESS FULL	EMP	1	14	14	00:00:00.01	4			

Predicate Information (identified by operation id):

4 - access("DEPTNO"="DEPTNO")

PLAN\_TABLE\_OUTPUT

filter("DEPTNO"="DEPTNO")

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			3	7
MERGE JOIN		SEMI	3	7
TABLE ACCESS	SCOTT.DEPT	BY INDEX ROWID	4	2
INDEX	SCOTT.PK_DEPT	FULL SCAN	4	1
SORT		UNIQUE	14	5
Access Predicates		DEPTNO=DEPTNO		
Filter Predicates		DEPTNO=DEPTNO		
TABLE ACCESS	SCOTT.EMP	FULL	14	4











Other XML

As we can see the inner join is not functionally equivalent to the semijoin because of the number of rows returned. The main difference between a normal inner join and a semijoin is that, with a semijoin, each record in the first set is returned only once, regardless of how many matches there are in the second set.



2.8. Task 9: Anti Joins

SQL WorksheetHistory








0,193 seconds

WorksheetQuery Builder

```
SELECT /* using in */ deptno, dname FROM scott.dept d
WHERE deptno NOT IN (select deptno from scott.emp e);
```

Script Output x



Task completed in 0,193 seconds

DEPTNO DNAME

-----

40 OPERATIONS

SELECT /\* using in \*/ deptno, dname FROM scott.dept d WHERE deptno NOT IN (select deptno from scott.emp e)

Plan hash value: 218628244

Id	Operation	Name	Starts	E-Rows	A-Rows	A-Time	Buffers	OMem	lMem	Used-Mem
0	SELECT STATEMENT		1		1	00:00:00.01	7			

PLAN\_TABLE\_OUTPUT

1	MERGE JOIN ANTI NA		1	1	1	00:00:00.01	7			
2	SORT JOIN		1	4	4	00:00:00.01	0	2048	2048	2048 (0)
3	TABLE ACCESS BY INDEX ROWID BATCHED	DEPT	1	4	4	00:00:00.01	2			
4	INDEX FULL SCAN	PK_DEPT	1	4	4	00:00:00.01	1			
* 5	SORT UNIQUE		5	14	3	00:00:00.01	0	2048	2048	2048 (0)
6	TABLE ACCESS FULL	EMP	1	14	14	00:00:00.01	5			

Predicate Information (identified by operation id):

PLAN\_TABLE\_OUTPUT

5 - access("DEPTNO"="DEPTNO")  
filter("DEPTNO"="DEPTNO")




























OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			1	7
MERGE JOIN		ANTI NA	1	7
SORT		JOIN	4	2
TABLE ACCESS	SCOTT.DEPT	BY INDEX ROWID...	4	2
INDEX	SCOTT.PK_DEPT	FULL SCAN	4	1
SORT		UNIQUE	14	5
Access Predicates		DEPTNO=DEPTNO		
Filter Predicates		DEPTNO=DEPTNO		
TABLE ACCESS	SCOTT.EMP	FULL	14	4

Other XML

2.9. Task 10: Prepare summary tableAnti-/Semi- Joins are applied depending



on the desired result(without repetition) and the selected Join type.

Join Access "A"	Join Access "B"	Nested Loop	Hash Join	Sort-Merge Join
Small Table	Small Table		 	 
Small Table	Indexed Small Table		 	 
Indexed Small	Indexed Small		  ok, but may be not as effective as others	
Small	Big		  ok, but hash table should be built for a smaller table	  ok if tables are sorted
Big	Big		 ok If the hash table can fit in memory	
Big	Indexed Big		 ok If the hash table can fit in memory	
Indexed Big	Indexed Big		 ok If the hash table can fit in memory	